

COMPUTATIONAL LEARNING IN DYNAMIC LOGICS

DAY 3: UPDATES ON NEURAL NETWORKS

Nina Gierasimczuk and Caleb Schultz Kisby

@NASSLLI, June 2025

Course Homepage:

<https://sites.google.com/view/nasslli25-learning-in-del>

PLAN FOR TODAY

- 1 Overview of Neural Networks
- 2 A Logic for Neural Network Inference
- 3 Neural Network Update in Dynamic Logic

PLAN FOR TODAY

- 1 Overview of Neural Networks
- 2 A Logic for Neural Network Inference
- 3 Neural Network Update in Dynamic Logic

A HISTORY OF NEURAL NETWORK LEARNING

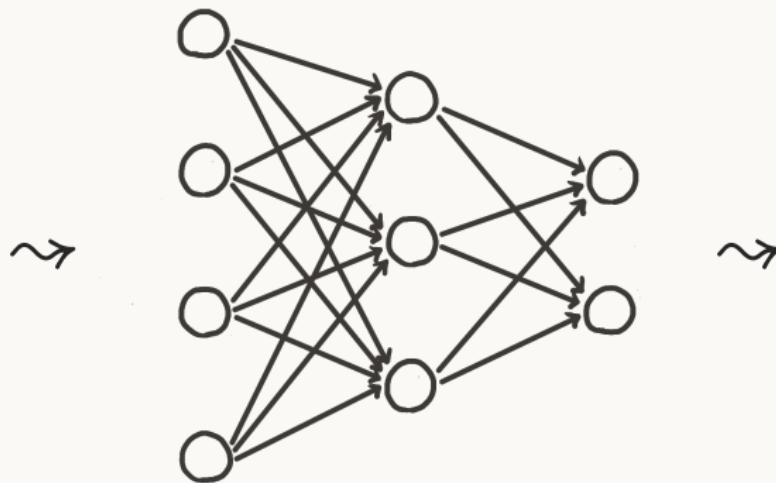
[keep it very general & current]

NEURAL NETWORKS IN MACHINE LEARNING

[keep it very general & current]

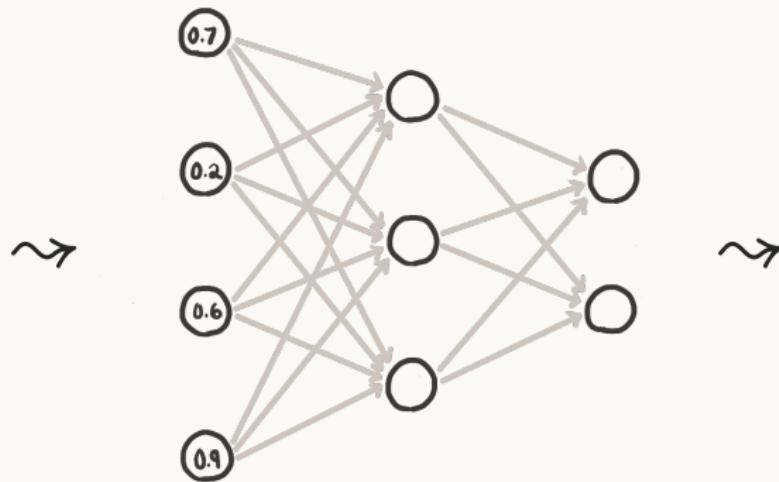
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:



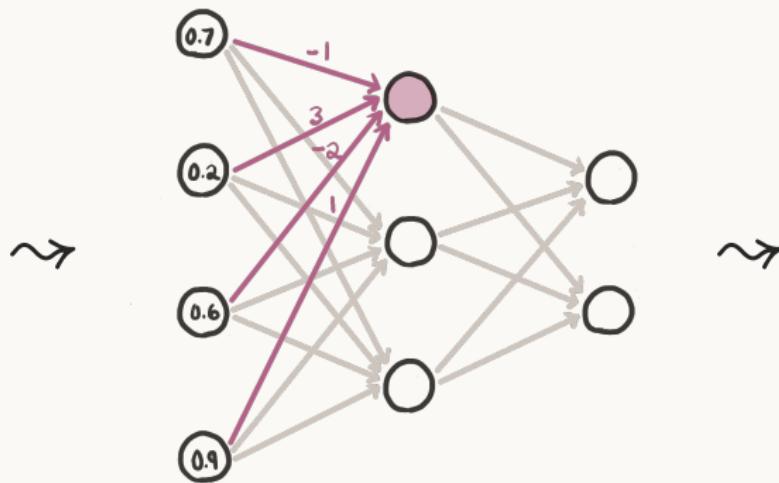
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:



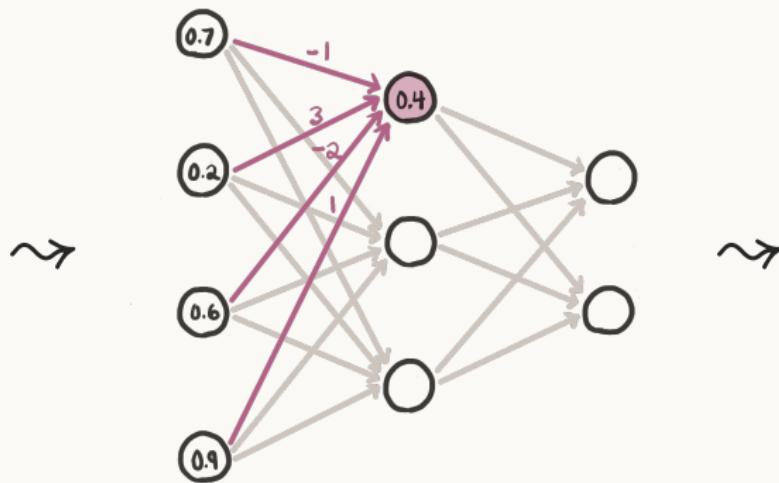
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:



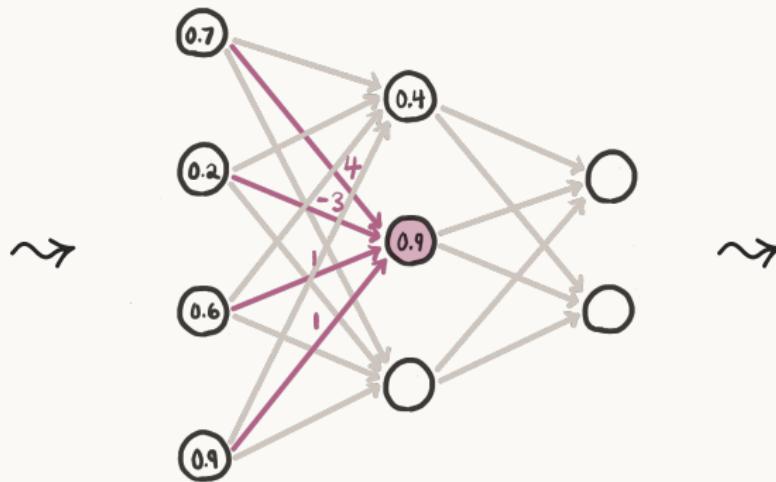
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:



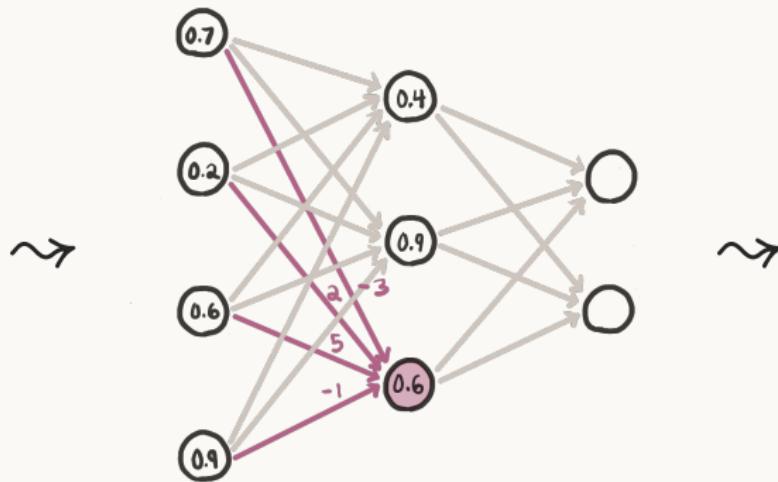
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:



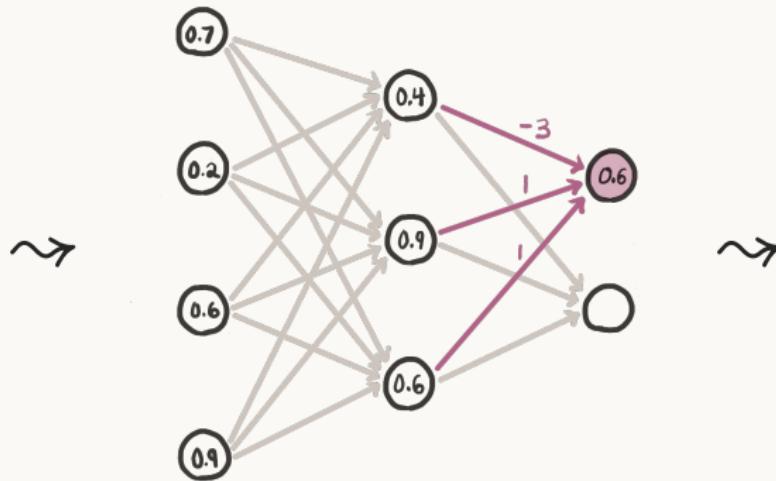
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:



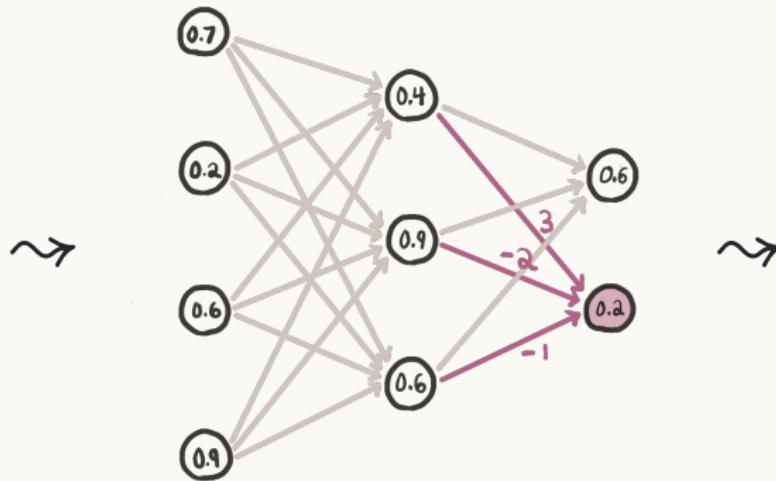
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:



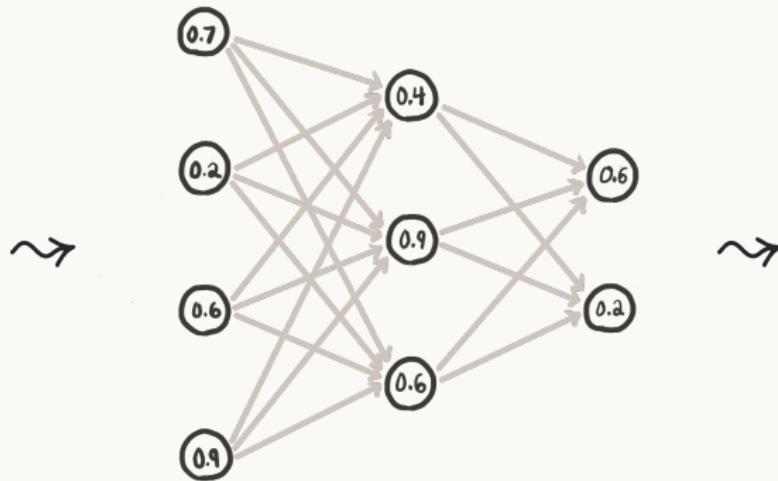
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:



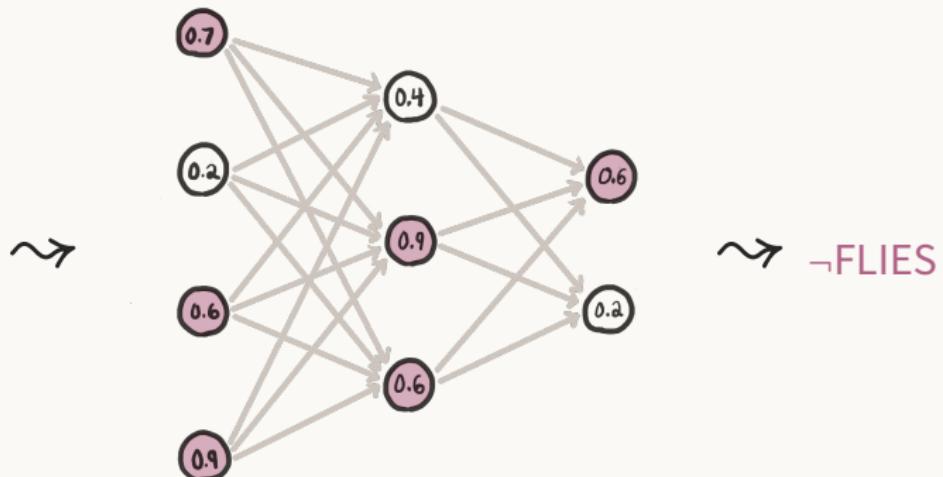
ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:

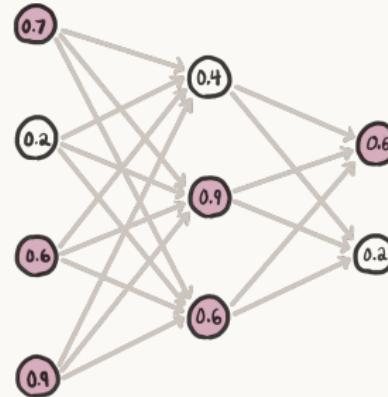
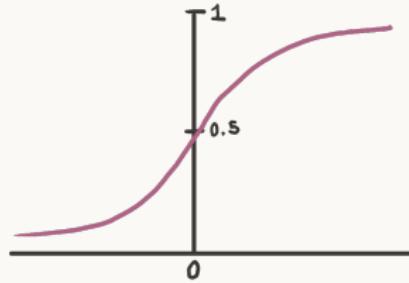


ARTIFICIAL NEURAL NETWORKS

- A neural network is just $\mathcal{N} = (N, E, W, A)$
 - neurons, edges, weights, activation function
- Neurons are successively activated by their predecessors:

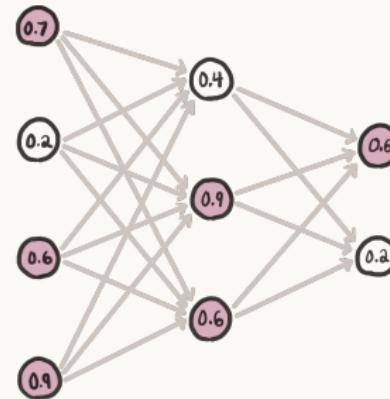
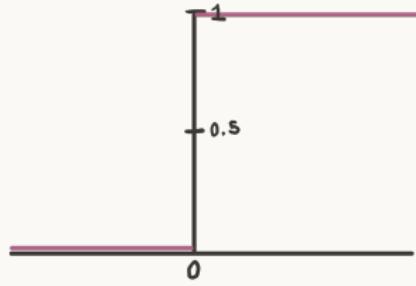


(BINARY) ARTIFICIAL NEURAL NETWORKS



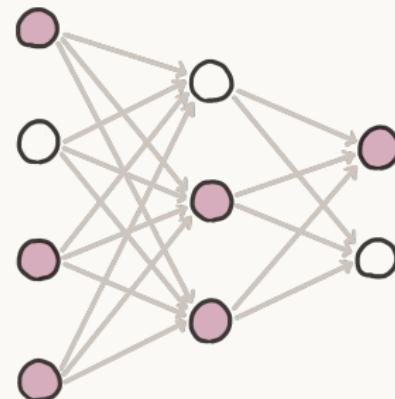
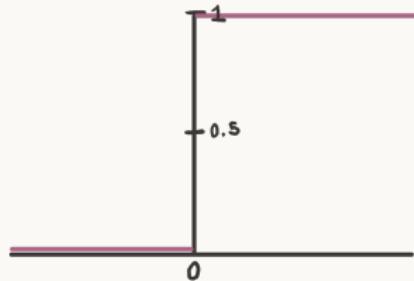
- We take the activation function A to be a binary step function
- This is a useful abstraction for connecting nets with logic, formal languages, and automata
- The net's activation patterns are just sets of neurons.

(BINARY) ARTIFICIAL NEURAL NETWORKS



- We take the activation function A to be a binary step function
- This is a useful abstraction for connecting nets with logic, formal languages, and automata
- The net's activation patterns are just sets of neurons.

(BINARY) ARTIFICIAL NEURAL NETWORKS



- We take the activation function A to be a binary step function
- This is a useful abstraction for connecting nets with logic, formal languages, and automata
- The net's activation patterns are just sets of neurons.

PLAN FOR TODAY

- 1 Overview of Neural Networks
- 2 A Logic for Neural Network Inference
- 3 Neural Network Update in Dynamic Logic

SYNTAX: LANGUAGE OF NEURAL NETWORK INFERENCE

Definition (Language of Epistemic Logic)

Take a countable set of propositions PROP.

$$\varphi := \top \mid p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \mathbf{A}\varphi \mid \langle \mathbf{C} \rangle \varphi$$

for all $p \in \text{PROP}$. The usual abbreviations are \vee , \rightarrow , and \mathbf{C} (dual to $\langle \mathbf{C} \rangle$)

SYNTAX: LANGUAGE OF NEURAL NETWORK INFERENCE

Definition (Language of Epistemic Logic)

Take a countable set of propositions PROP.

$$\varphi := \top \mid p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \mathbf{A}\varphi \mid \langle \mathbf{C} \rangle \varphi$$

for all $p \in \text{PROP}$. The usual abbreviations are \vee , \rightarrow , and \mathbf{C} (dual to $\langle \mathbf{C} \rangle$)

- Notice that we're giving the semantics in terms of the \diamond -variant $\langle \mathbf{C} \rangle$

SYNTAX: LANGUAGE OF NEURAL NETWORK INFERENCE

Definition (Language of Epistemic Logic)

Take a countable set of propositions PROP.

$$\varphi := \top \mid p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \mathbf{A}\varphi \mid \langle \mathbf{C} \rangle \varphi$$

for all $p \in \text{PROP}$. The usual abbreviations are \vee , \rightarrow , and \mathbf{C} (dual to $\langle \mathbf{C} \rangle$)

- Notice that we're giving the semantics in terms of the \diamond -variant $\langle \mathbf{C} \rangle$
- Just like before, we will interpret $\langle \mathbf{C} \rangle \varphi$ in a model, at a world.

SYNTAX: LANGUAGE OF NEURAL NETWORK INFERENCE

Definition (Language of Epistemic Logic)

Take a countable set of propositions PROP.

$$\varphi := \top \mid p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \mathbf{A}\varphi \mid \langle \mathbf{C} \rangle \varphi$$

for all $p \in \text{PROP}$. The usual abbreviations are \vee , \rightarrow , and \mathbf{C} (dual to $\langle \mathbf{C} \rangle$)

- Notice that we're giving the semantics in terms of the \diamond -variant $\langle \mathbf{C} \rangle$
- Just like before, we will interpret $\langle \mathbf{C} \rangle \varphi$ in a model, at a world.
- **The intended interpretation:**

$\langle \mathbf{C} \rangle \varphi$ holds in a net, at a neuron w if w is activated by input φ .

SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- A binary neural network is just $\mathcal{N} = (N, E, W, A)$

SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- A binary neural network is just $\mathcal{N} = (N, E, W, A)$
- Each choice of E, W, A specifies a transition function from one activation pattern $S \subseteq N$ to the next

SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- A binary neural network is just $\mathcal{N} = (N, E, W, A)$
- Each choice of E, W, A specifies a transition function from one activation pattern $S \subseteq N$ to the next
- Given initial state S_0 , $F_{S_0} : \wp(N) \rightarrow \wp(N)$ is given by

$$F_{S_0}(S) = S_0 \cup \{w \mid A\left(\sum_{u \in \text{preds}(w)} W(u, w) \cdot \chi_S(u)\right) = 1\}$$

“the set of all nodes w activated by their immediate predecessors u ”

SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- A binary neural network is just $\mathcal{N} = (N, E, W, A)$
- Each choice of E, W, A specifies a transition function from one activation pattern $S \subseteq N$ to the next
- Given initial state S_0 , $F_{S_0} : \wp(N) \rightarrow \wp(N)$ is given by

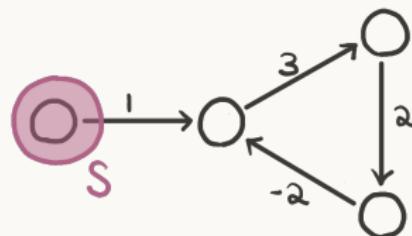
$$F_{S_0}(S) = S_0 \cup \{w \mid A\left(\sum_{u \in \text{preds}(w)} W(u, w) \cdot \chi_S(u)\right) = 1\}$$

“the set of all nodes w activated by their immediate predecessors u ”

- $\chi_S(u) = 1$ iff $u \in S$ indicates whether u was activated previously

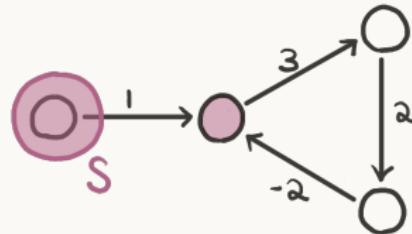
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



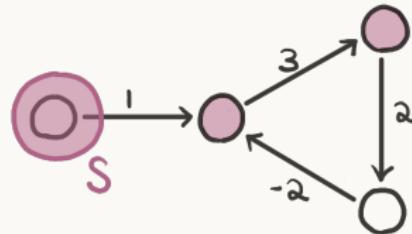
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



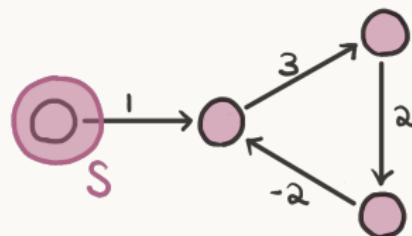
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



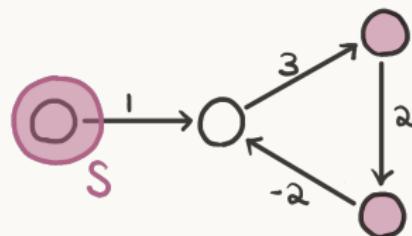
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



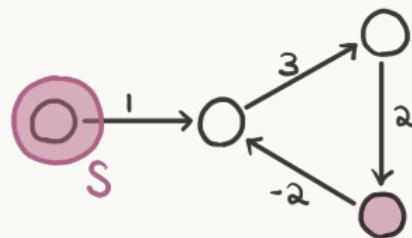
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



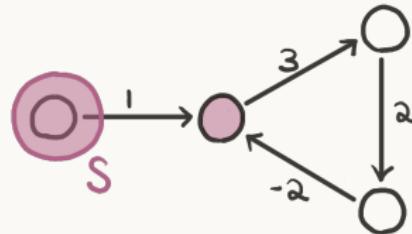
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



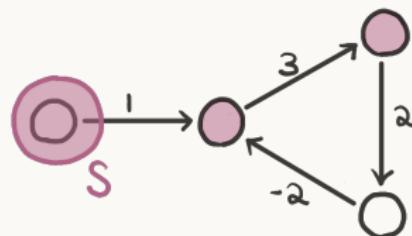
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



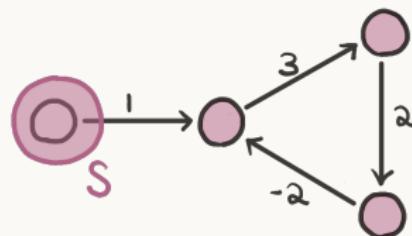
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



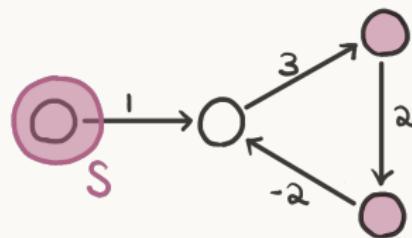
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



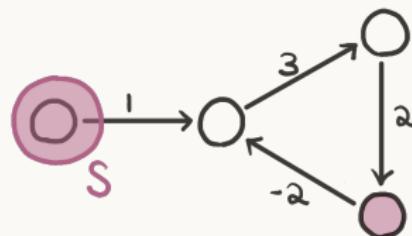
SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

- Notice that the activated nodes could have oscillatory behavior!
- But we only want nets that have a unique “answer” for each input



SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

Postulate

We assume for all $S_0 \subseteq N$, F_{S_0} repeatedly applied to S_0 ,

$$S_0, F_{S_0}(S_0), F_{S_0}(F_{S_0}(S_0)), \dots, F_{S_0}^k(S_0), \dots$$

eventually stabilizes to a unique activation pattern.

SEMANTICS: NEURAL NETWORK CLOSURE OPERATOR

Postulate

We assume for all $S_0 \subseteq N$, F_{S_0} repeatedly applied to S_0 ,

$$S_0, F_{S_0}(S_0), F_{S_0}(F_{S_0}(S_0)), \dots, F_{S_0}^k(S_0), \dots$$

eventually stabilizes to a unique activation pattern.

Definition

Let $\text{Clos} : \wp(N) \rightarrow \wp(N)$ be the function that produces this stable activation pattern.

SEMANTICS: FORMAL DEFINITION

Definition (Neural Network Semantics)

Given a binary neural network model $\mathcal{N} = (N, E, W, A, V)$, where

$V : Prop \rightarrow \wp(N)$, and a neuron (“world”) $w \in N$:

$$\mathcal{N}, w \models p \quad \text{iff} \quad w \in V(p) \text{ for each } p \in Prop$$

$$\mathcal{N}, w \models \neg\varphi \quad \text{iff} \quad \text{not } \mathcal{N}, w \models \varphi$$

$$\mathcal{N}, w \models \varphi \wedge \psi \quad \text{iff} \quad \mathcal{N}, w \models \varphi \text{ and } \mathcal{N}, w \models \psi$$

$$\mathcal{N}, w \models A\varphi \quad \text{iff} \quad \text{for all } w \in N \text{ whatsoever, } \mathcal{N}, w \models \varphi$$

$$\mathcal{N}, w \models \langle \mathbf{C} \rangle \varphi \quad \text{iff} \quad w \in \text{Clos}(\llbracket \varphi \rrbracket)$$

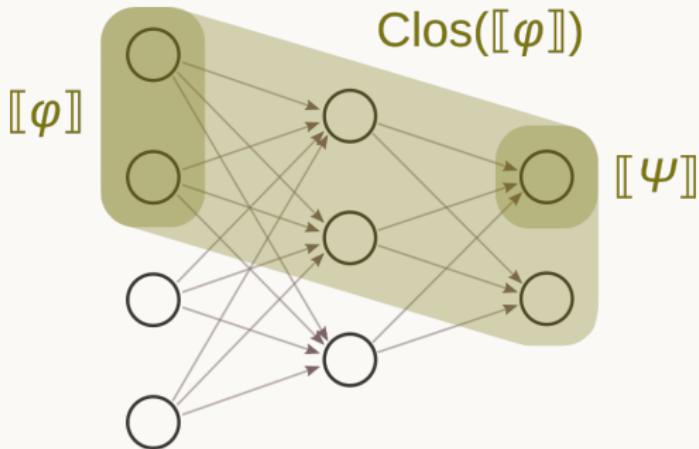
and dually:

$$\mathcal{N}, w \models \mathbf{C}\varphi \quad \text{iff} \quad w \in (\text{Clos}(\llbracket \varphi \rrbracket))^{\mathbf{C}}$$

where $\llbracket \varphi \rrbracket = \{u \mid \mathcal{N}, u \models \varphi\}$ is the set of worlds where φ holds (the set of neurons that are active for φ)

EXPRESSING NEURAL NETWORK INFERENCE

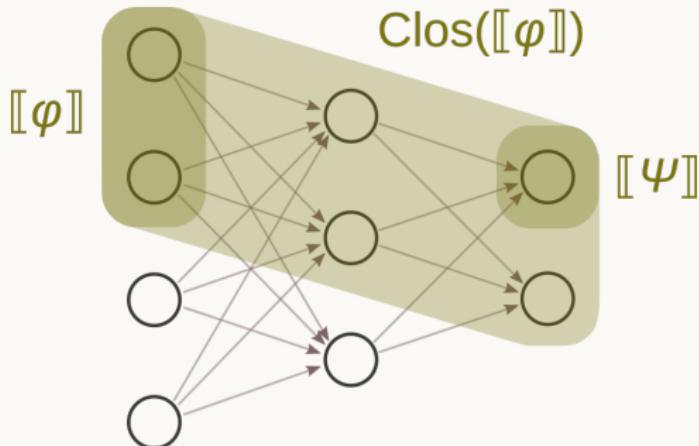
- The **C** modality gives information about the net's answer to an input



The net satisfies $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$ iff The net satisfies $\mathbf{A}(\psi \rightarrow \langle \mathbf{C} \rangle(\varphi))$

EXPRESSING NEURAL NETWORK INFERENCE

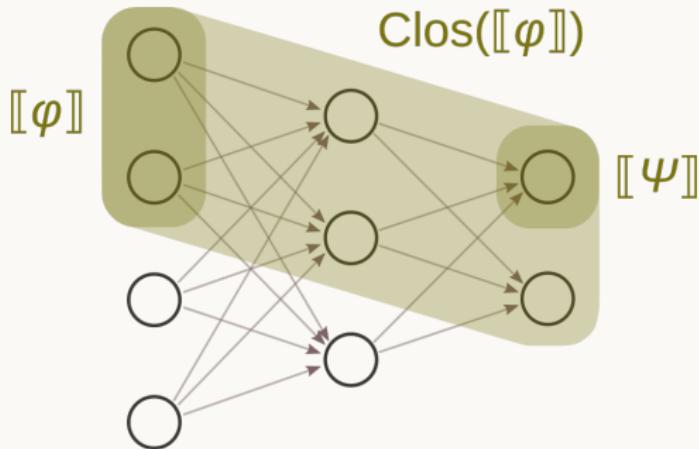
- The **C** modality gives information about the net's answer to an input



The net satisfies $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$ iff The net satisfies $\mathbf{A}(\psi \rightarrow \langle \mathbf{C} \rangle(\varphi))$
iff $\text{Clos}(\llbracket \varphi \rrbracket) \supseteq \llbracket \psi \rrbracket$

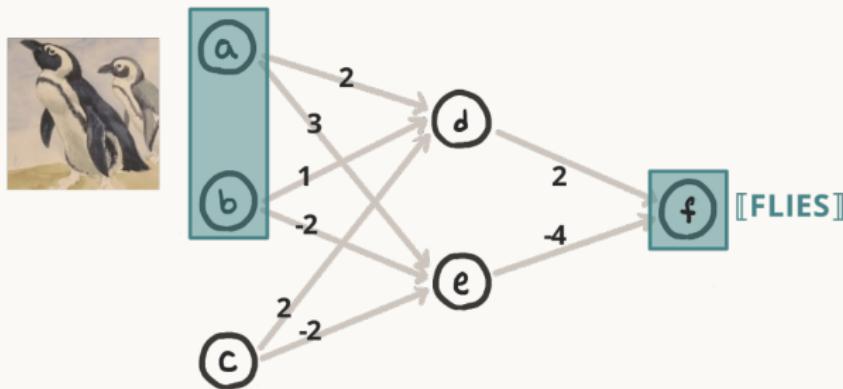
EXPRESSING NEURAL NETWORK INFERENCE

- The **C** modality gives information about the net's answer to an input



The net satisfies $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$ iff The net satisfies $\mathbf{A}(\psi \rightarrow \langle \mathbf{C} \rangle(\varphi))$
iff $\text{Clos}([\varphi]) \supseteq [\psi]$
iff **The net classifies φ as ψ**

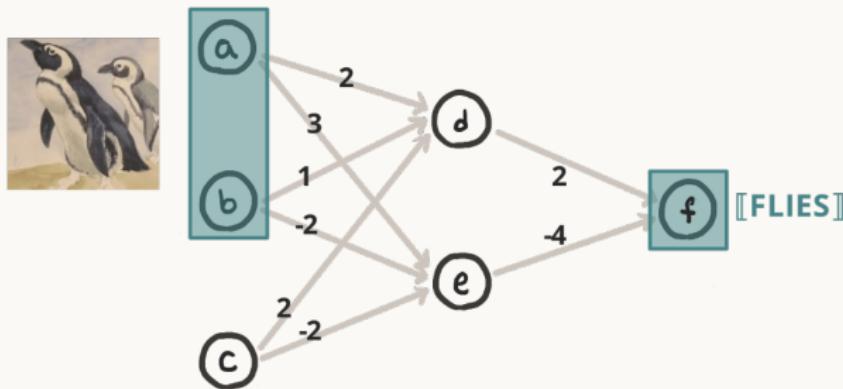
EXAMPLE: EXPRESSING NEURAL NETWORK INFERENCE



- In the exercises, we will ask you will to show

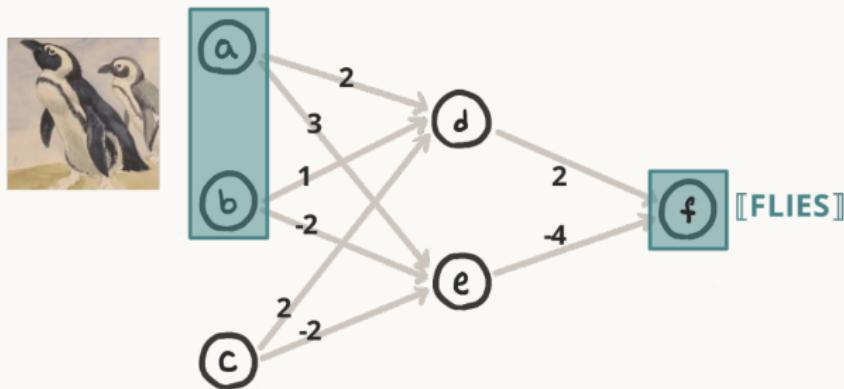
$$\mathcal{N} \not\models A(C(PENGUIN) \rightarrow FLIES)$$

EXAMPLE: EXPRESSING NEURAL NETWORK INFERENCE



- In the exercises, we will ask you will to show
$$\mathcal{N} \not\models \mathbf{A}(\mathbf{C}(\text{PENGUIN}) \rightarrow \text{FLIES})$$
- This means the net does not classify penguins as flying

EXAMPLE: EXPRESSING NEURAL NETWORK INFERENCE



- In the exercises, we will ask you will to show

$$\mathcal{N} \not\models \mathbf{A}(\mathbf{C}(\text{PENGUIN}) \rightarrow \text{FLIES})$$

- This means the net does not classify penguins as flying
- Yet, if we take 〔BIRD〕 = {a, b, c},

$$\mathcal{N} \models \mathbf{A}(\mathbf{C}(\text{BIRD}) \rightarrow \text{FLIES})$$

ADDITIONAL COMMENTS ON THIS LOGIC

The net satisfies $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$ iff The net satisfies $\mathbf{A}(\psi \rightarrow \langle \mathbf{C} \rangle(\varphi))$
iff $\text{Clos}([\![\varphi]\!]) \supseteq [\![\psi]\!]$
iff **The net classifies φ as ψ**

- What does this remind you of?

ADDITIONAL COMMENTS ON THIS LOGIC

The net satisfies $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$ iff The net satisfies $\mathbf{A}(\psi \rightarrow \langle \mathbf{C} \rangle(\varphi))$
iff $\text{Clos}([\![\varphi]\!]) \supseteq [\![\psi]\!]$
iff **The net classifies φ as ψ**

- What does this remind you of?
 - $\text{best}_\leq([\![\varphi]\!]) \subseteq [\![\psi]\!]$

ADDITIONAL COMMENTS ON THIS LOGIC

The net satisfies $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$ iff The net satisfies $\mathbf{A}(\psi \rightarrow \langle \mathbf{C} \rangle(\varphi))$
iff $\text{Clos}([\![\varphi]\!]) \supseteq [\![\psi]\!]$
iff **The net classifies φ as ψ**

- What does this remind you of?
 - $\text{best}_\leq([\![\varphi]\!]) \subseteq [\![\psi]\!]$
 - $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$, taken as a conditional, behaves exactly like $B^\varphi \psi$

ADDITIONAL COMMENTS ON THIS LOGIC

The net satisfies $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$ iff The net satisfies $\mathbf{A}(\psi \rightarrow \langle \mathbf{C} \rangle(\varphi))$
iff $\text{Clos}([\![\varphi]\!]) \supseteq [\![\psi]\!]$
iff **The net classifies φ as ψ**

- What does this remind you of?
 - $\text{best}_\leq([\![\varphi]\!]) \subseteq [\![\psi]\!]$
 - $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$, taken as a conditional, behaves exactly like $B^\varphi \psi$
 - You can think of this as the net's conditional belief

ADDITIONAL COMMENTS ON THIS LOGIC

The net satisfies $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$ iff The net satisfies $\mathbf{A}(\psi \rightarrow \langle \mathbf{C} \rangle(\varphi))$
iff $\text{Clos}([\![\varphi]\!]) \supseteq [\![\psi]\!]$
iff **The net classifies φ as ψ**

- What does this remind you of?
 - $\text{best}_{\leq}([\![\varphi]\!]) \subseteq [\![\psi]\!]$
 - $\mathbf{A}(\mathbf{C}(\varphi) \rightarrow \psi)$, taken as a conditional, behaves exactly like $B^{\varphi}\psi$
 - You can think of this as the net's conditional belief
- Interpreting **C** on its own is less clear...

PLAN FOR TODAY

- 1 Overview of Neural Networks
- 2 A Logic for Neural Network Inference
- 3 Neural Network Update in Dynamic Logic

UPDATES ON NEURAL NETWORKS

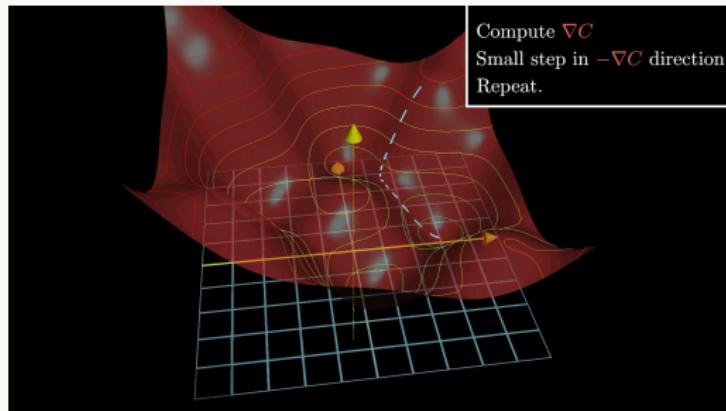
- **Unsupervised Updates**
 - The network learns from data that is unlabeled (no expected answer or classification)
 - Each update softly increases the net's preference for the input
 - Hebb's rule, Oja's rule, & competitive learning rule
- **Supervised Updates**
 - The network learns from labeled data with an expected answer
 - Each update softly increases the net's accuracy on a function
 - Backpropagation rule & delta learning rule

UPDATES ON NEURAL NETWORKS

- **Unsupervised Updates**
 - The network learns from data that is unlabeled (no expected answer or classification)
 - Each update softly increases the net's preference for the input
 - **Hebb's rule**, Oja's rule, & competitive learning rule
- **Supervised Updates**
 - The network learns from labeled data with an expected answer
 - Each update softly increases the net's accuracy on a function
 - **Backpropagation rule** & delta learning rule

BACKPROPAGATION RULE

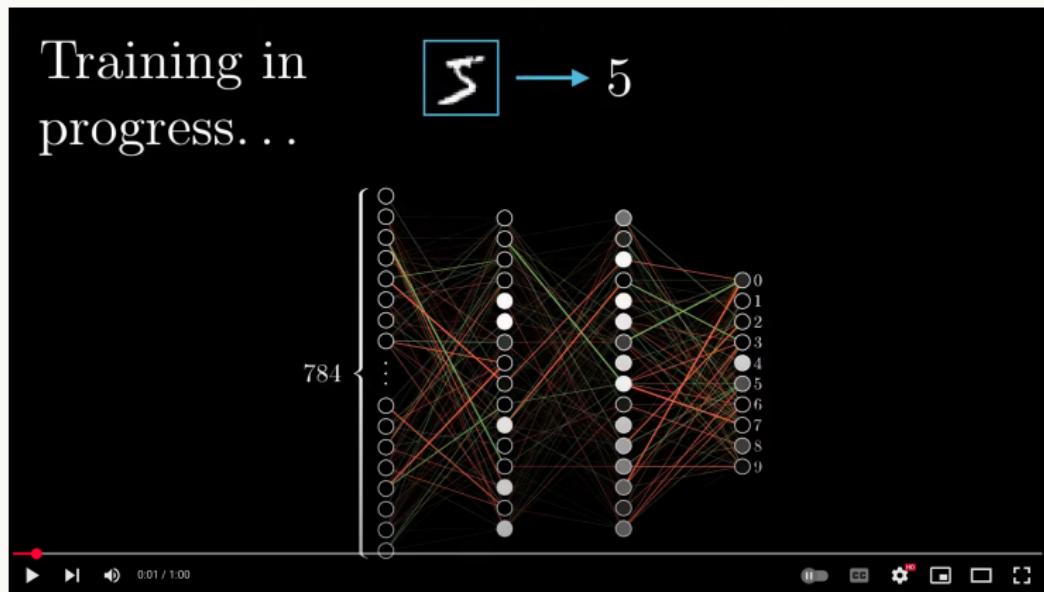
- Backpropagation is the most widely used neural network update rule
- **Main idea:** Backprop implements gradient descent on a net's weights



- Given an input \vec{x} with label y , the neural network gives its answer y' to \vec{x} , and each weight of the net is adjusted according to its contribution to the error (difference between y' and y).

BACKPROPAGATION RULE

<https://www.youtube.com/watch?v=cANqroNVdI8>



ADDITIONAL COMMENTS ON NEURAL NETWORK UPDATES

- What if we could have a complete characterization of Backprop, like we did for public announcement, L_{EX}, and MINI?
 - That would be wonderful!

ADDITIONAL COMMENTS ON NEURAL NETWORK UPDATES

- What if we could have a complete characterization of Backprop, like we did for public announcement, L_{EX}, and MINI?
 - That would be wonderful!
 - Unfortunately, this is still an open problem

ADDITIONAL COMMENTS ON NEURAL NETWORK UPDATES

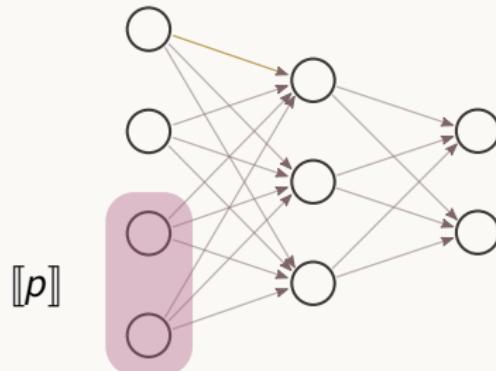
- What if we could have a complete characterization of Backprop, like we did for public announcement, L_{EX}, and MINI?
 - That would be wonderful!
 - Unfortunately, this is still an open problem
- **Proof of concept:** Can we do this for any neural network update at all?

ADDITIONAL COMMENTS ON NEURAL NETWORK UPDATES

- What if we could have a complete characterization of Backprop, like we did for public announcement, L_{EX}, and MINI?
 - That would be wonderful!
 - Unfortunately, this is still an open problem
- **Proof of concept:** Can we do this for any neural network update at all?
 - Let's consider the simplest possible one: **Hebbian learning**

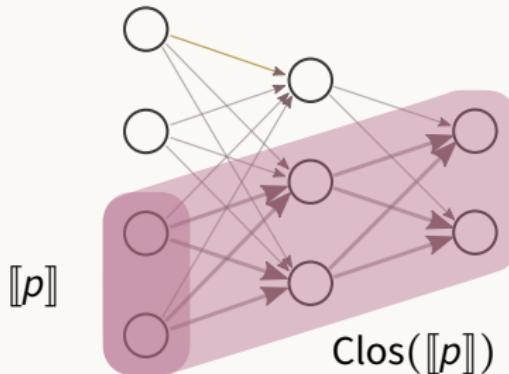
HEBBIAN UPDATE RULE

Neurons that fire together wire together



HEBBIAN UPDATE RULE

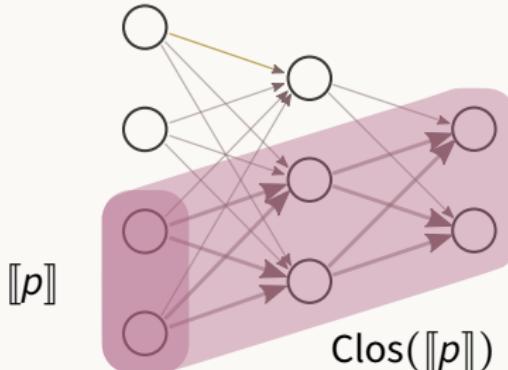
Neurons that fire together wire together



- Each edge involved in the activation is “bumped up” by a fixed learning rate $\eta \geq 0$

HEBBIAN UPDATE RULE

Neurons that fire together wire together



- Each edge involved in the activation is “bumped up” by a fixed learning rate $\eta \geq 0$
- Formally: $\text{Hebb}(N, \llbracket \varphi \rrbracket) = (N, E, W', A)$, where $W'(u, w) = W(u, w) + \eta \cdot \chi_{\text{Clos}(\llbracket \varphi \rrbracket)}(u) \cdot \chi_{\text{Clos}(\llbracket \varphi \rrbracket)}(w)$

HEBBIAN UPDATE RULE: EXAMPLES

- If nobody ever told you that penguins don't fly, how could you come to believe they don't?

HEBBIAN UPDATE RULE: EXAMPLES

- If nobody ever told you that penguins don't fly, how could you come to believe they don't?
 - Observe animals with similar features that **don't** fly

HEBBIAN UPDATE RULE: EXAMPLES

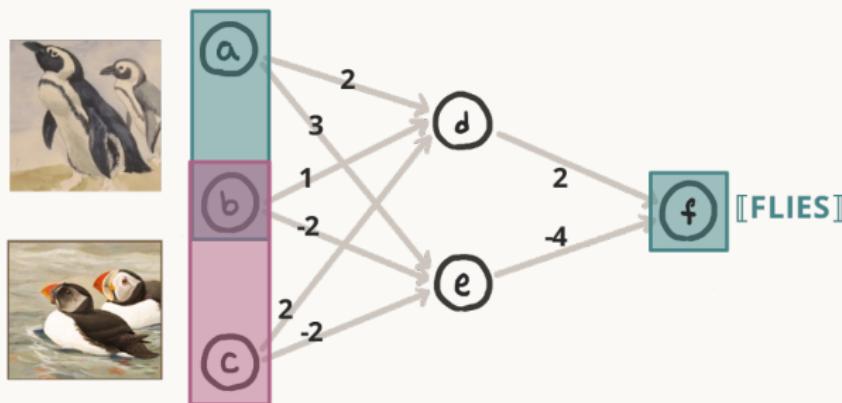
- If nobody ever told you that penguins don't fly, how could you come to believe they don't?
 - Observe animals with similar features that **don't** fly
- Now imagine you believe penguins don't fly. What could cause you to change your mind?

HEBBIAN UPDATE RULE: EXAMPLES

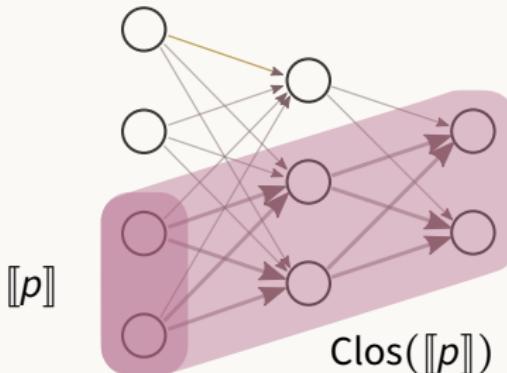
- If nobody ever told you that penguins don't fly, how could you come to believe they don't?
 - Observe animals with similar features that **don't** fly
- Now imagine you believe penguins don't fly. What could cause you to change your mind?
 - Observe animals with similar features that **do** fly

HEBBIAN UPDATE RULE: EXAMPLES

- If nobody ever told you that penguins don't fly, how could you come to believe they don't?
 - Observe animals with similar features that **don't** fly
- Now imagine you believe penguins don't fly. What could cause you to change your mind?
 - Observe animals with similar features that **do** fly

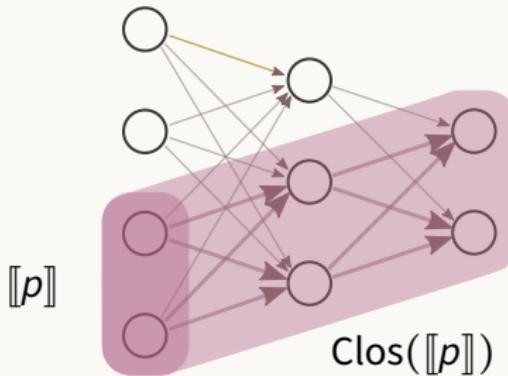


HEBB*: “FIXED-POINT” HEBBIAN UPDATE



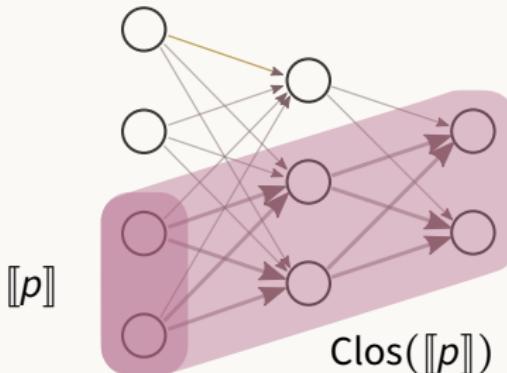
- If we repeatedly apply $\text{Hebb}((, N), S)$, eventually these weights will saturate (they will not inhibit any incoming activations)

HEBB*: “FIXED-POINT” HEBBIAN UPDATE



- If we repeatedly apply $\text{Hebb}((, N), S)$, eventually these weights will saturate (they will not inhibit any incoming activations)
- Let iter be the number of iterations needed to reach this fixed point

HEBB*: “FIXED-POINT” HEBBIAN UPDATE



- If we repeatedly apply $\text{Hebb}((N), S)$, eventually these weights will saturate (they will not inhibit any incoming activations)
- Let iter be the number of iterations needed to reach this fixed point
- Let $\text{Hebb}^*(N, S) = (N, E, W', A)$, where

$$W'(u, w) = W(u, w) + \text{iter} \cdot \eta \cdot \chi_{\text{Clos}(\llbracket \varphi \rrbracket)}(u) \cdot \chi_{\text{Clos}(\llbracket \varphi \rrbracket)}(w)$$

NEURAL NETWORK UPDATES IN DYNAMIC LOGIC

- We can use the DEL trick to give semantics using neural network updates

Definition (Neural Network Semantics)

Let \mathcal{N} be a binary neural network model, $w \in N$, and let

$\mathcal{U} : \mathbf{Net} \rightarrow \mathcal{L} \rightarrow \mathbf{Net}$ be any unsupervised update:

$$\mathcal{N}, w \vDash [\varphi]\psi \quad \text{iff} \quad \text{Update}(N, [[\varphi]]), w \vDash \psi$$

For Hebbian updates in particular:

$$\mathcal{N}, w \vDash [\varphi]_{Hebb}\psi \quad \text{iff} \quad \text{Hebb}(N, [[\varphi]]), w \vDash \psi$$

$$\mathcal{N}, w \vDash [\varphi]_{Hebb^*}\psi \quad \text{iff} \quad \text{Hebb}^*(N, [[\varphi]]), w \vDash \psi$$

REDUCTION LAWS FOR HEBB*

The following formulas are valid over neural network models:

REDUCTION LAWS FOR HEBB^{*}

The following formulas are valid over neural network models:

Hebb^{*}-Propositions $[\varphi]_{\text{Hebb}^*} p \leftrightarrow p$

REDUCTION LAWS FOR HEBB^{*}

The following formulas are valid over neural network models:

Hebb^{*}-Propositions $[\varphi]_{\text{Hebb}^*} p \leftrightarrow p$

Hebb^{*}-Negation $[\varphi]_{\text{Hebb}^*} \neg \psi \leftrightarrow \neg [\varphi]_{\text{Hebb}^*} \psi$

REDUCTION LAWS FOR HEBB*

The following formulas are valid over neural network models:

Hebb*-Propositions $[\varphi]_{\text{Hebb}^*} p \leftrightarrow p$

Hebb*-Negation $[\varphi]_{\text{Hebb}^*} \neg \psi \leftrightarrow \neg [\varphi]_{\text{Hebb}^*} \psi$

Hebb*-Conjunction $[\varphi]_{\text{Hebb}^*} (\psi \wedge \theta) \leftrightarrow [\varphi]_{\text{Hebb}^*} \psi \wedge [\varphi]_{\text{Hebb}^*} \theta$

REDUCTION LAWS FOR HEBB*

The following formulas are valid over neural network models:

Hebb*-Propositions $[\varphi]_{\text{Hebb}^*} p \leftrightarrow p$

Hebb*-Negation $[\varphi]_{\text{Hebb}^*} \neg \psi \leftrightarrow \neg [\varphi]_{\text{Hebb}^*} \psi$

Hebb*-Conjunction $[\varphi]_{\text{Hebb}^*} (\psi \wedge \theta) \leftrightarrow [\varphi]_{\text{Hebb}^*} \psi \wedge [\varphi]_{\text{Hebb}^*} \theta$

Hebb*-Diamond $[\varphi]_{\text{Hebb}^*} \diamond \psi \leftrightarrow \diamond [\varphi]_{\text{Hebb}^*} \psi$

REDUCTION LAWS FOR HEBB*

The following formulas are valid over neural network models:

Hebb*-Propositions $[\varphi]_{\text{Hebb}^*} p \leftrightarrow p$

Hebb*-Negation $[\varphi]_{\text{Hebb}^*} \neg \psi \leftrightarrow \neg [\varphi]_{\text{Hebb}^*} \psi$

Hebb*-Conjunction $[\varphi]_{\text{Hebb}^*} (\psi \wedge \theta) \leftrightarrow [\varphi]_{\text{Hebb}^*} \psi \wedge [\varphi]_{\text{Hebb}^*} \theta$

Hebb*-Diamond $[\varphi]_{\text{Hebb}^*} \diamond \psi \leftrightarrow \diamond [\varphi]_{\text{Hebb}^*} \psi$

Hebb*-Closure $[\varphi]_{\text{Hebb}^*} \langle \mathbf{C} \rangle \psi \leftrightarrow$
 $\langle \mathbf{C} \rangle ([\varphi]_{\text{Hebb}^*} \psi \vee (\langle \mathbf{C} \rangle \varphi \wedge \diamond (\langle \mathbf{C} \rangle \varphi \wedge \langle \mathbf{C} \rangle [\varphi]_{\text{Hebb}^*} \psi)))$

REDUCTION LAWS FOR HEBB*

The following formulas are valid over neural network models:

Hebb*-Propositions $[\varphi]_{\text{Hebb}^*} p \leftrightarrow p$

Hebb*-Negation $[\varphi]_{\text{Hebb}^*} \neg \psi \leftrightarrow \neg [\varphi]_{\text{Hebb}^*} \psi$

Hebb*-Conjunction $[\varphi]_{\text{Hebb}^*} (\psi \wedge \theta) \leftrightarrow [\varphi]_{\text{Hebb}^*} \psi \wedge [\varphi]_{\text{Hebb}^*} \theta$

Hebb*-Diamond $[\varphi]_{\text{Hebb}^*} \diamond \psi \leftrightarrow \diamond [\varphi]_{\text{Hebb}^*} \psi$

Hebb*-Closure $[\varphi]_{\text{Hebb}^*} \langle \mathbf{C} \rangle \psi \leftrightarrow$
 $\langle \mathbf{C} \rangle ([\varphi]_{\text{Hebb}^*} \psi \vee (\langle \mathbf{C} \rangle \varphi \wedge \diamond (\langle \mathbf{C} \rangle \varphi \wedge \langle \mathbf{C} \rangle [\varphi]_{\text{Hebb}^*} \psi)))$

Note: We didn't define \diamond for neural networks; here are its semantics:

$\mathcal{N}, w \models \diamond \varphi$ iff there is an E -path from some $u \in \llbracket \varphi \rrbracket$ to w .

TO LEARN MORE CHECK OUT...

Neural Network Models of Conditionals: An Introduction

Humanistic

Department of Philosophy, University of Bristol,
Hannes Leitgeb@bristol.ac.uk

114

This "instant mines-style" article gives a brief survey of neural network models of conditionals. After some introductory remarks on the nature of neural networks and conditionals, we turn to the notion of an *intensional dynamical system* as a unifying concept in the logical investigation of dynamic systems in general, and of neural networks in particular. We explain how conditionals get represented by intensional dynamical systems, which logical systems these conditionals obey, and what the main open problem is in this area.

Keywords: Conditionals, Neural networks, Dynamical systems, Nonmonotonic logic.

1 Introduction

Neural networks are abstract models of brain structures capable of adapting to new information. The learning abilities of artificial neural networks have given rise to successful computer implementations of various cognitive tasks, from the recognition of

Formal logics are the *precursors of common sense*. Logic deals with formal systems of reasoning; in particular, inductive logic studies formal systems of reasoning towards plausible but uncautious conclusions. As evidence accumulates, the degree to which it supports a hypothesis, as measured by the logic, should tend to indicate that the hypothesis is likely to be true.

two areas developed in opposition to each other: neural networks are quantitative dynamic systems, while logical reasoners must be symbolic systems; networks are described by mathematical equations, whereas logic is subject to intuitive statements about how we ought to reason; neural networks have been studied by scientists, whilst the "problem of induction" is regarded as belonging to the humanities.

**(a) Neural Network
Models of Conditionals:
An Introduction by
Leitgeb**

The Warm, Fuzzy, and Continuously Artificial Intelligence Model 48

What Do Hebbian Learners Learn? Reduction Axioms for Iterated Hebbian Learning

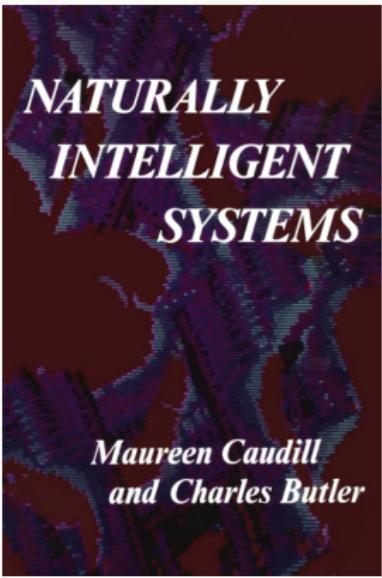
Caleb Schultz Kishy¹, Saúl A. Blanco¹, Lawrence S. Mass²
¹Department of Computer Science, Indiana University
²Department of Mathematics, Indiana University
Bloomington, IN 47408, USA
{ckishy, sblanco, lmass}@indiana.edu

Abstract
The central question this theory **Soundsness**: What are sounds? Answer: “Any sound operation can map the same input to different outputs” [1].
Completeness: What are the complete network operations? This is equivalent to asking “Can we build a network that can do everything?” [2]. Can we build a classical model?
We refer the reader to the landmark of Ávila-García 2022, which shows a comparison in a wide class of neural-style generative models that we consider the most important.

Proposes a set N . Active neurons in N activate new neurons until eventual saturation. Each EC neuron has one or more active neurons.

and symbolic systems, have long sought increased interdisciplinarity and symbiosis. But, are there benefits to it and is there a danger in the overemphasis of one over the other? In this paper, we argue that, while the field of neural networks has been instrumental in advancing the state-of-the-art in many domains, the field of symbolic AI has emerged as a complementary mode of inquiry in cognitive science and symbolic systems. We argue that the two fields can benefit from each other by creating a symbiotic relationship that can lead to more robust and general AI systems. In this paper, we will introduce the basic concepts of neural networks and symbolic systems and show how they can be used together to create a symbiotic system.

(b) Reduction Axioms for Iterated Hebbian Learning by Schultz Kisby, Blanco, & Moss



(c) **Naturally Intelligent Systems** by Caudill & Butler

END OF LECTURE 3

Thank you!